

# Polynomial Regression

## Importing the Libraries

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

## Importing the dataset

```
In [2]: dataset = pd.read_csv('Position_Salaries.csv')
X = dataset.iloc[:, 1:2].values
y = dataset.iloc[:, 2].values
```

## Splitting the dataset into the Training set and Test set

*Since we need the entire dataset for training (for this particular problem), we will not split the data into training and test sets*

```
In [ ]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 0)
```

## Fitting Linear Regression to the dataset

```
In [3]: from sklearn.linear_model import LinearRegression
lin_reg = LinearRegression()
lin_reg.fit(X, y)
```

```
Out[3]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

## Fitting Polynomial Regression to the dataset

```
In [4]: from sklearn.preprocessing import PolynomialFeatures
poly_reg = PolynomialFeatures(degree = 4)
X_poly = poly_reg.fit_transform(X)
poly_reg.fit(X_poly, y)
lin_reg_2 = LinearRegression()
lin_reg_2.fit(X_poly, y)
```


```
Out[4]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

*As we can see above, the higher degrees of the expression leads to better fit. This however, can also culminate into a over-fitted model.*

*Please ensure that the regressor is not overfitted to suit the training data*


## Visualising the Linear Regression results

```
In [5]: plt.scatter(X, y, color = 'red')
plt.plot(X, lin_reg.predict(X), color = 'blue')
plt.title('Truth or Bluff (Linear Regression)')
plt.xlabel('Position level')
plt.ylabel('Salary')
plt.show()
```



## Visualising the Polynomial Regression results

```
In [6]: plt.scatter(X, y, color = 'red')
plt.plot(X, lin_reg_2.predict(poly_reg.fit_transform(X)), color = 'blue')
plt.title('Truth or Bluff (Polynomial Regression)')
plt.xlabel('Position level')
plt.ylabel('Salary')
plt.show()
```



## Predicting a new result with Linear Regression

```
In [7]: lin_reg.predict(6.5)
```

```
Out[7]: array([330378.78787879])
```

## Predicting a new result with Polynomial Regression

```
In [8]: lin_reg_2.predict(poly_reg.fit_transform(6.5))
```

```
Out[8]: array([158862.4526516])
```